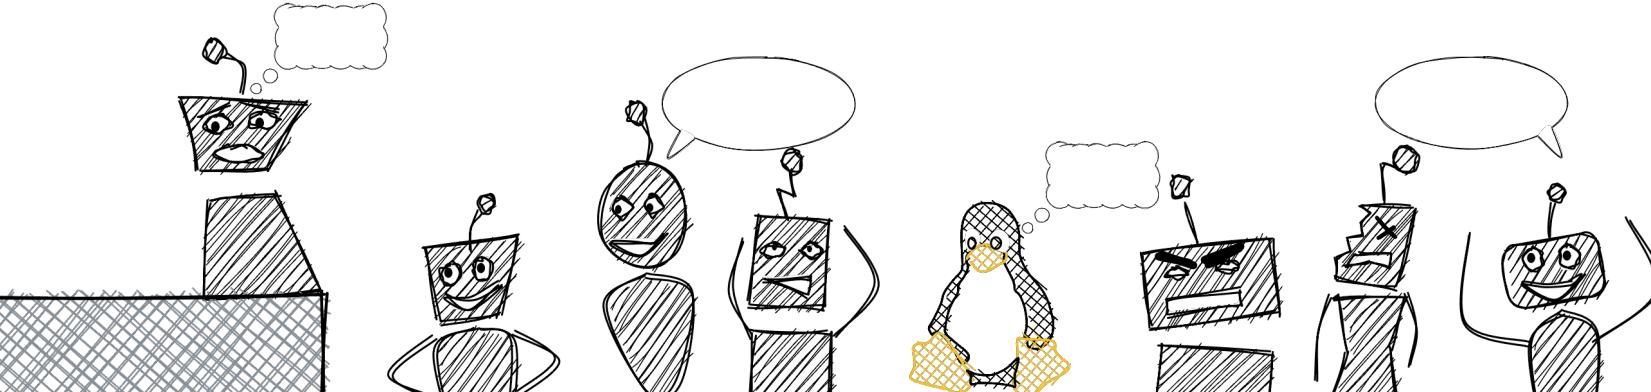


The 4C's of Cloud Native Security



Valentin Hristev
Staff Site Reliability Engineer



The 4C's of Cloud Native Security



Container and Image security



Runtime security



Kubernetes pod security



Kubernetes network security



Behavioral Analysis



Supply Chain Security





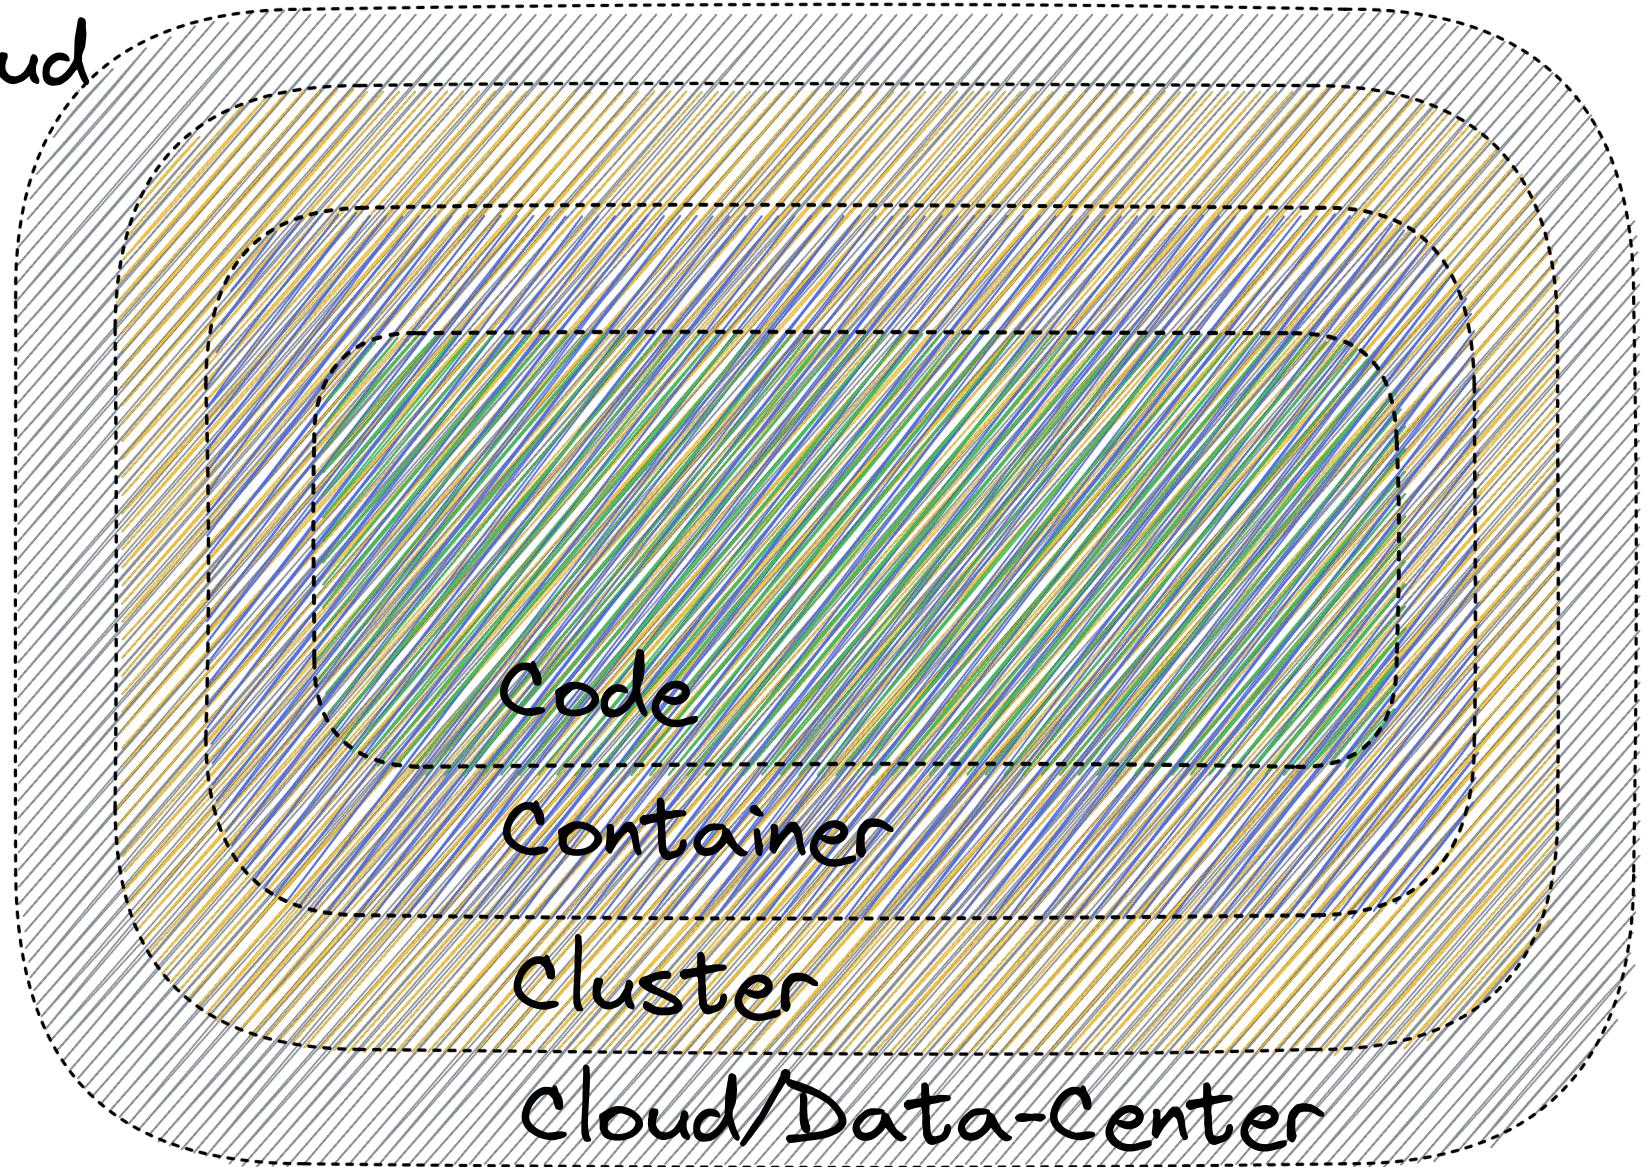
The 4C's of Cloud Native Security



Security is a complex multi-dimensional **beast**, and it must be contained on different levels.



The 4C's of Cloud Native Security



Cloud / Datacenter security

- The 6 Layers of Google Data Center Security
 - Layer 1: Signage and Fencing
 - Layer 2: Secure Perimeter
 - Layer 3: Building Access
 - Layer 4: Security Operations Center
 - Layer 5: Data Center Floor
 - Layer 6: Secure Hard Drive Destruction

<https://www.youtube.com/watch?v=kd33UVZhnAA&t=37s>



Container and Image security

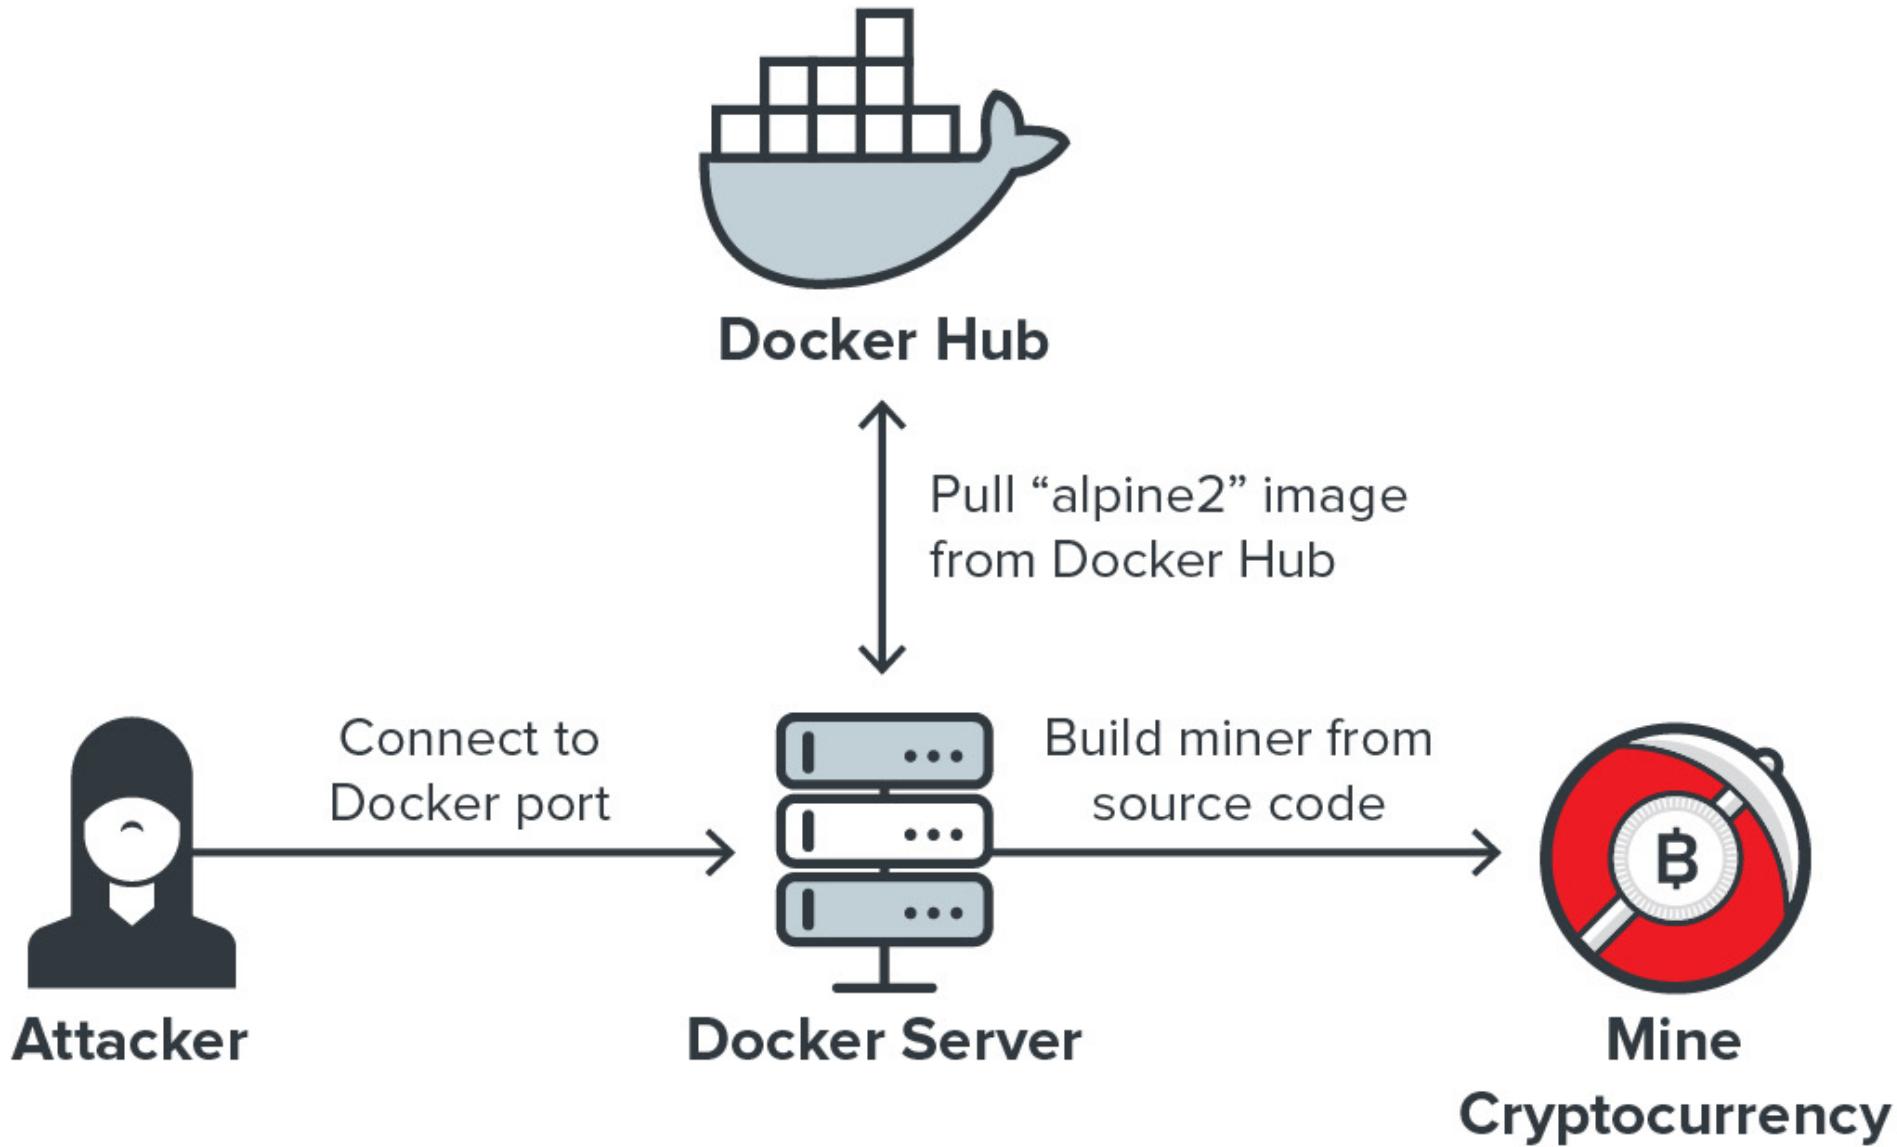


Public registries

- Can we trust them?
 - Yes, you can trust me



"By pushing malicious images to a Docker Hub registry and pulling it from the victim's system, hackers were able to mine 544.74 Monero, which equals \$90,000."

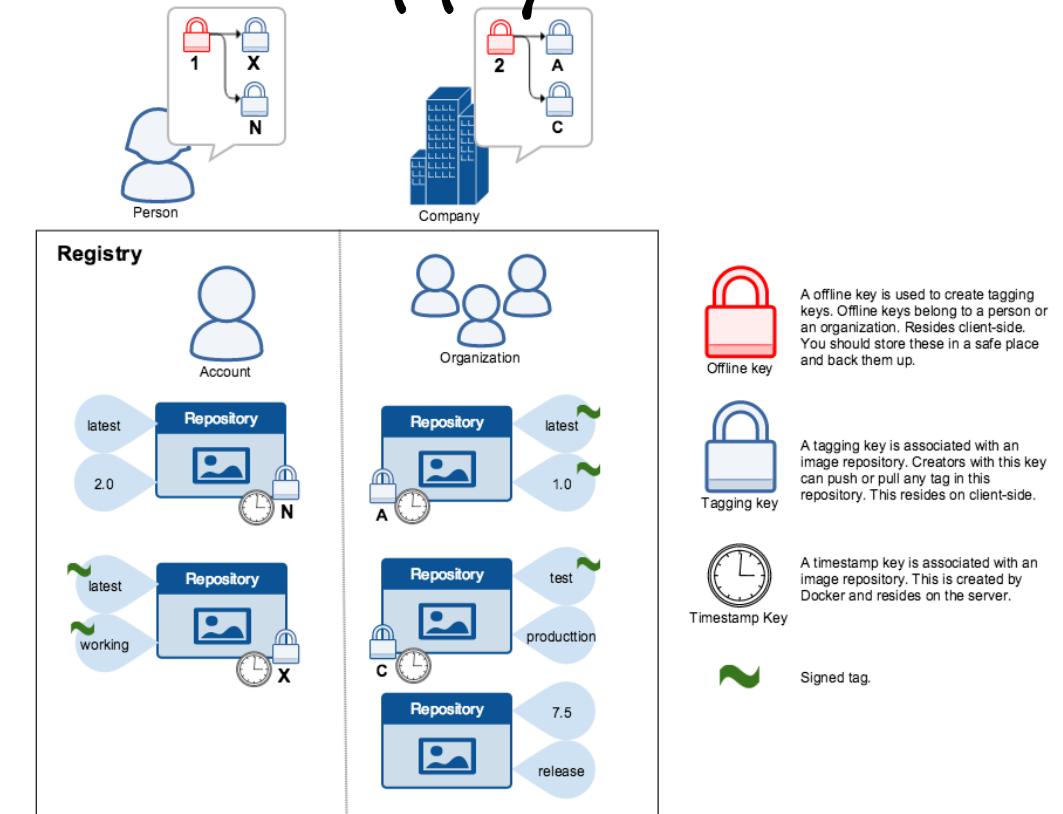


Build images best practices

- Multi-stage build
 - Drop build tools/libs from the final image
- Run as a non-root user
 - Don't run as root just because it's easy
 - I need port below 1024 that's why I need root
- Old images like 'Ubuntu:12.04'
- tag images
 - Are tags mutable or immutable?
- source code at the last step (use cache)

Sign and validate images

- Image scanning included in the supply chain
 - Trivy
- Sign image
 - Gnu Privacy Guard (GPG)
 - Docker Content Trust (DCT)
- Timestamp key





Runtime security



Open container Initiative

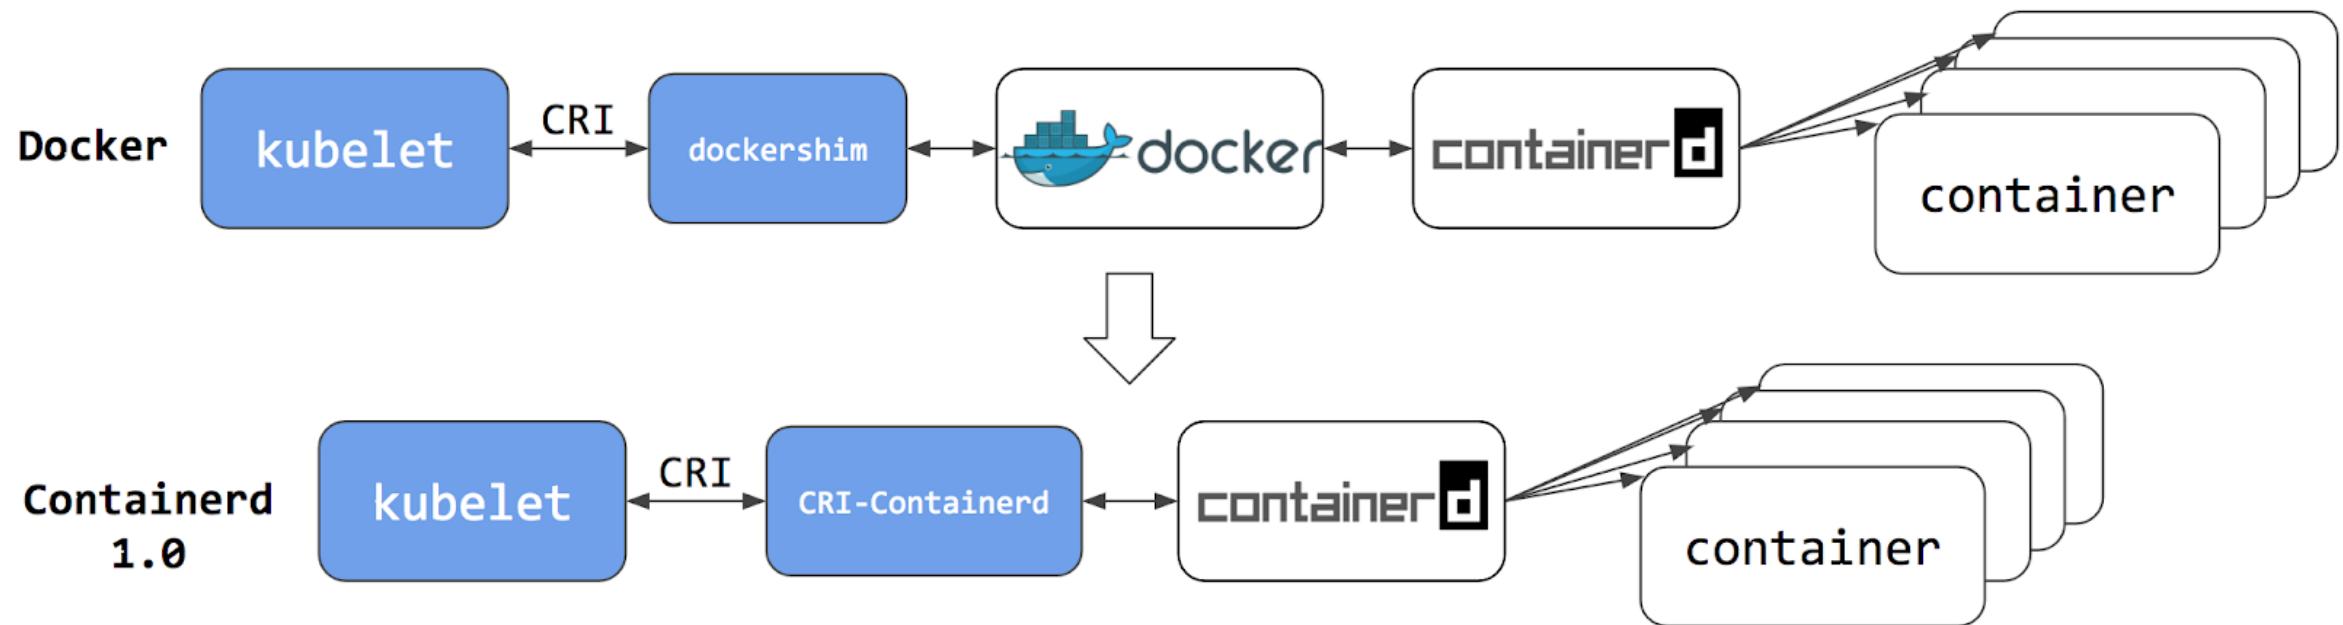
- Runtime specification
 - Specify execution env and container lifecycle
- Image specification
 - Download, unpack and run it specification
- Distribution specification
 - Standardize container image distribution based on Dock Registry HTTP API v2

Lets run our image

- Container Runtime Interface (CRI)
- OCI runners - low level container exec
 - runc - written in **Go** (Docker is using it)
 - crun - redhat written in **C**
- CRI shims - pull and runs images per OCI
 - cri-dockerd
 - containerd
 - cri-o
- Container Engines
 - Docker - all in one building, publishing, running ...

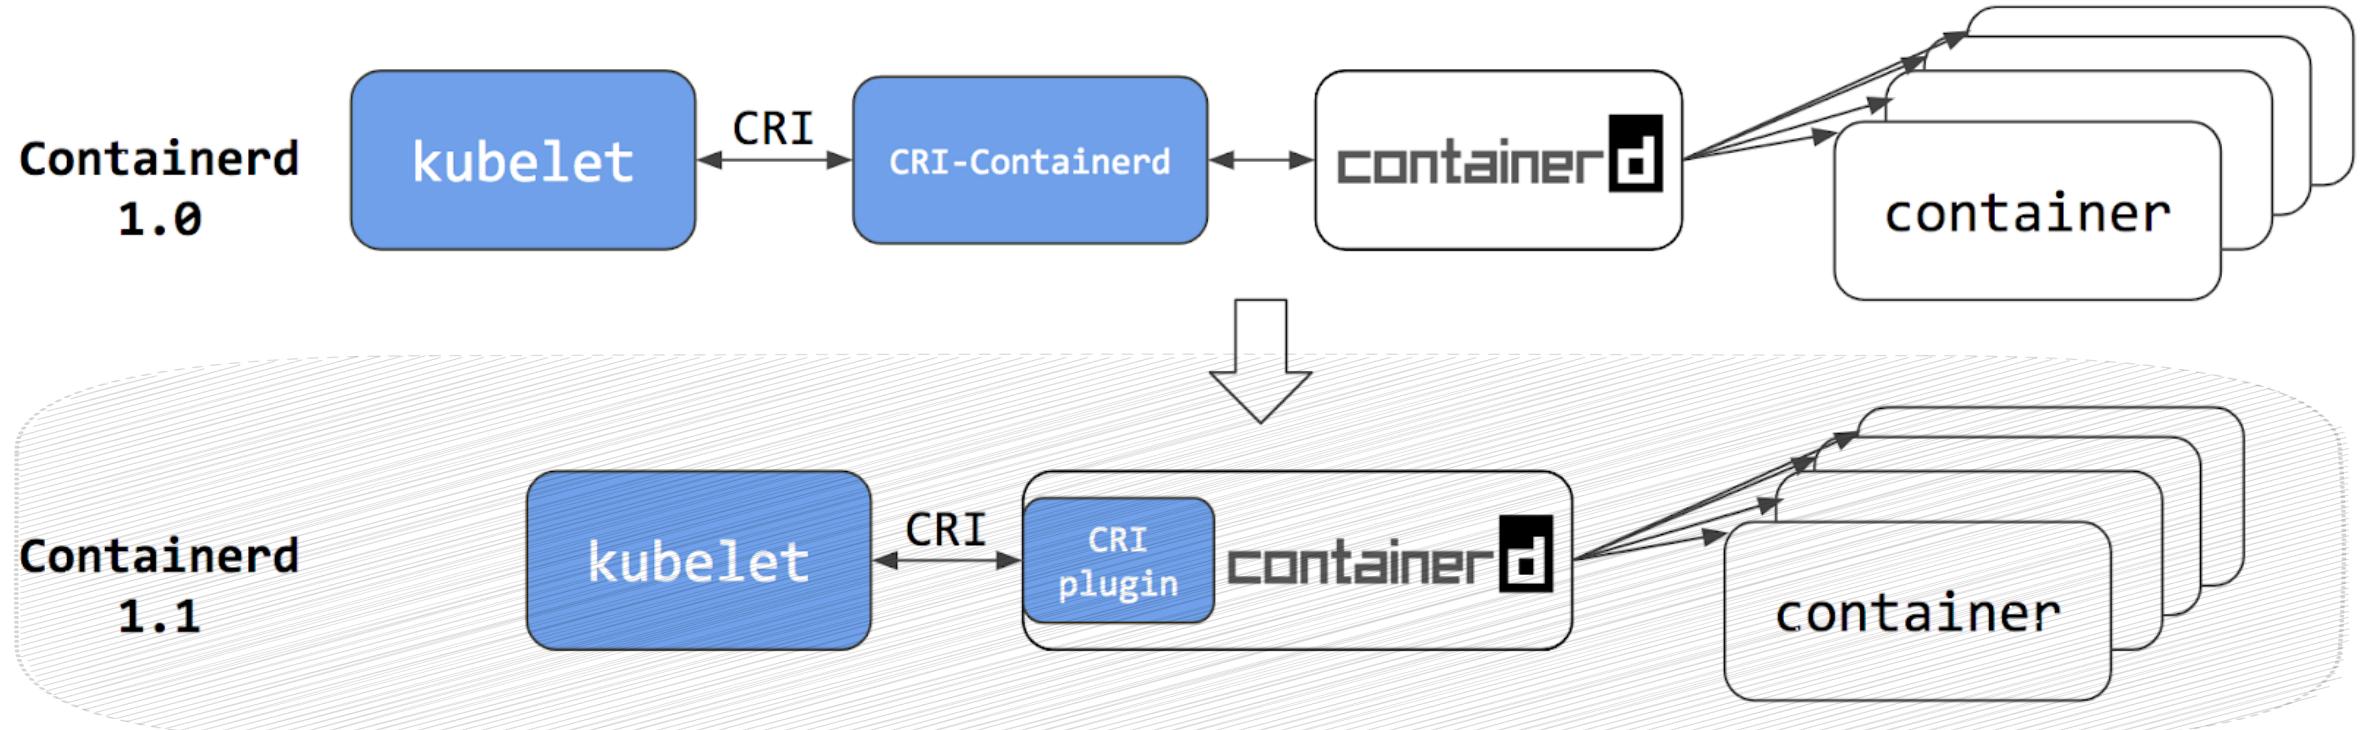
Lets run our image

- Container Runtime Interface (CRI)



Lets run our image

- Container Runtime Interface (CRI)

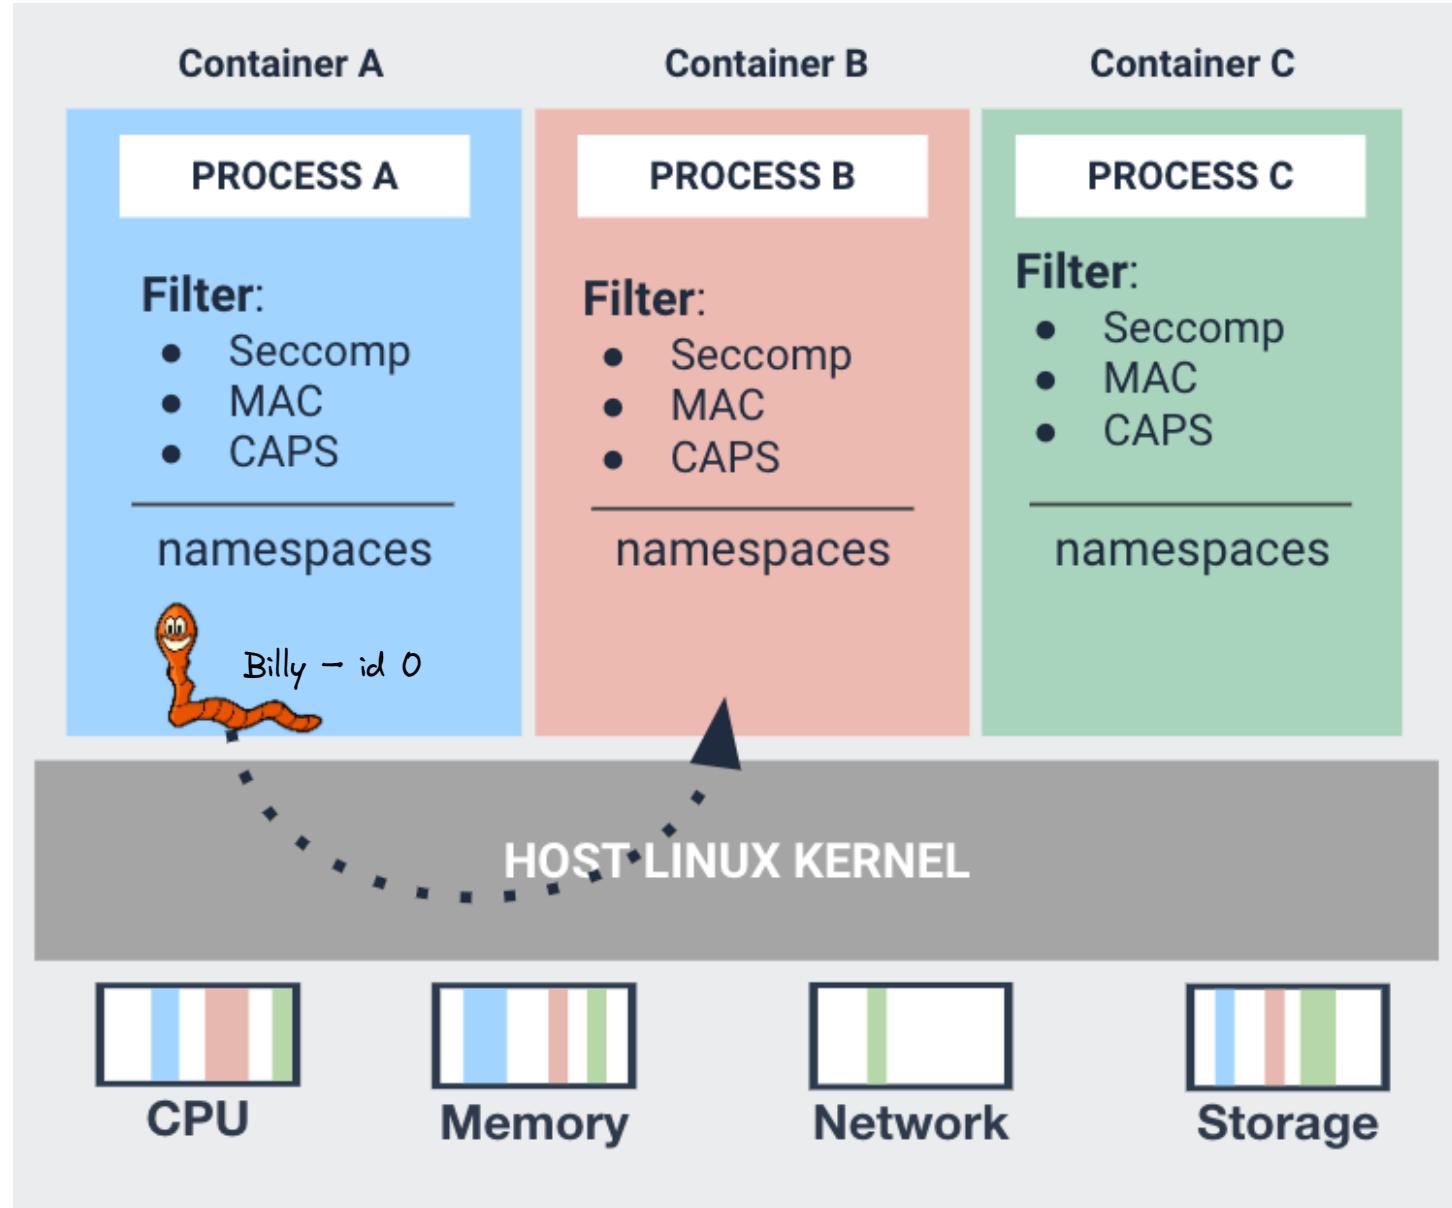


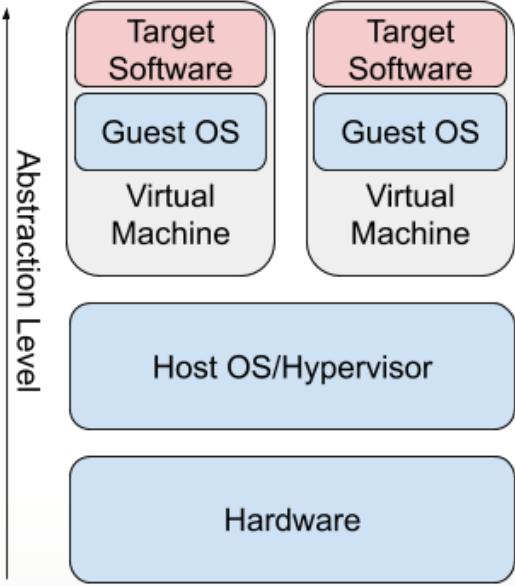
The OCI is the standard
not Docker

Docker is OCI compliant

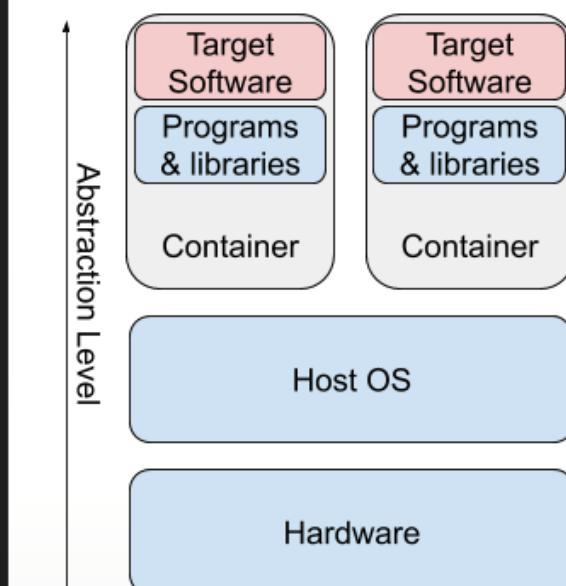
Container vs VM vs Lightvisor

- Linux containers - There is **NO** such thing as a '**container**'
- Known the limits of the namespace isolation
- Do not run different customers on the same physical host

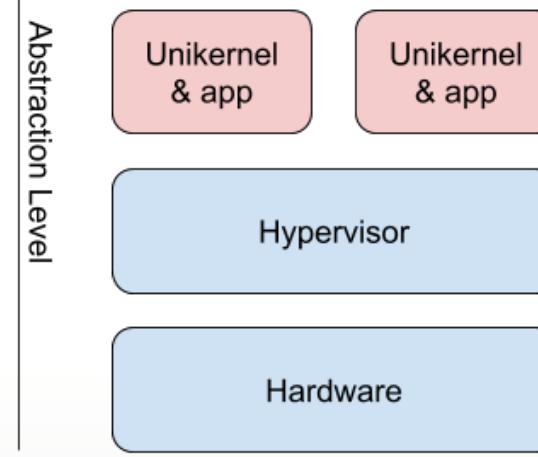




Virtual Machine
(Multi tenant)



Container

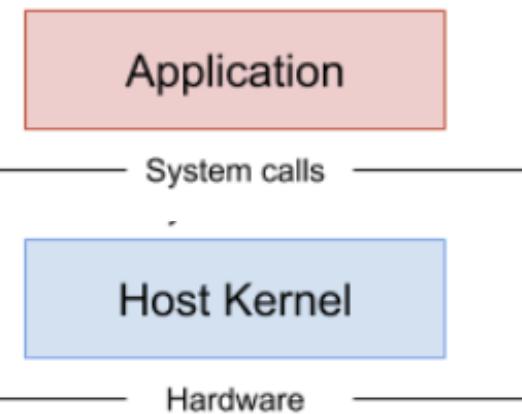


Unikernel
Kata / gVisor

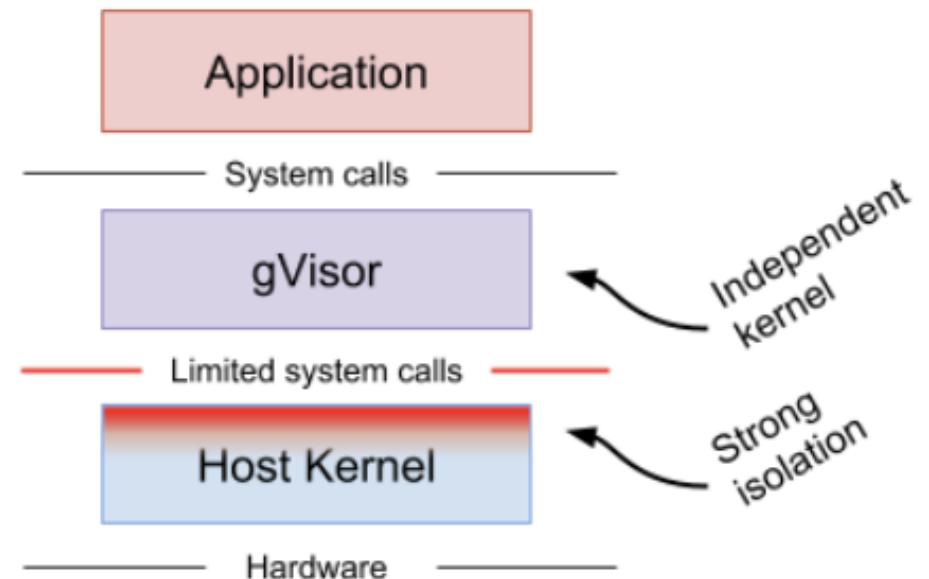
```
~$ sudo crictl inspect c86389dc7fa48 | grep -iA 20 namespaces
namespaces: [
    {
        "type": "ipc",
        "path": "/proc/10413/ns/ipc"
    },
    {
        "path": "/proc/10413/ns/uts"
    },
    {
        "path": "/proc/10413/ns/net"
    }
]

$ ps -ef | grep 10413
nobody 10413 10375 0 10:06 ? 00:00:00 runsc-sandbox --
root=/run/containerd/runsc/k8s.io --
log=/run/containerd/io.containerd.runtime.v2.task/k8s.io/63df904cf22f2b44adf5
a79c69b31f80aabe54aa9f53ae60351b87114fce416a/log.json --log-format=json --
log-fd=3 boot --
bundle=/run/containerd/io.containerd.runtime.v2.task/k8s.io/63df904cf22f2b44a
df5a79c69b31f80aabe54aa9f53ae60351b87114fce416a -- controller-fd=4 --
mounts-fd=5 --spec-fd=6 --start-sync-fd=7 --io-fds=8 --io-fds=9 --stdio-fds=10
--stdio-fds=11 --stdio-fds=12 --cpu-num 2 --user-log-fd 13
63df904cf22f2b44adf5a79c69b31f80aabe54aa9f53ae60351b87114fce416a

nobody 10484 10413 0 10:06 ? 00:00:00 [exe]
```

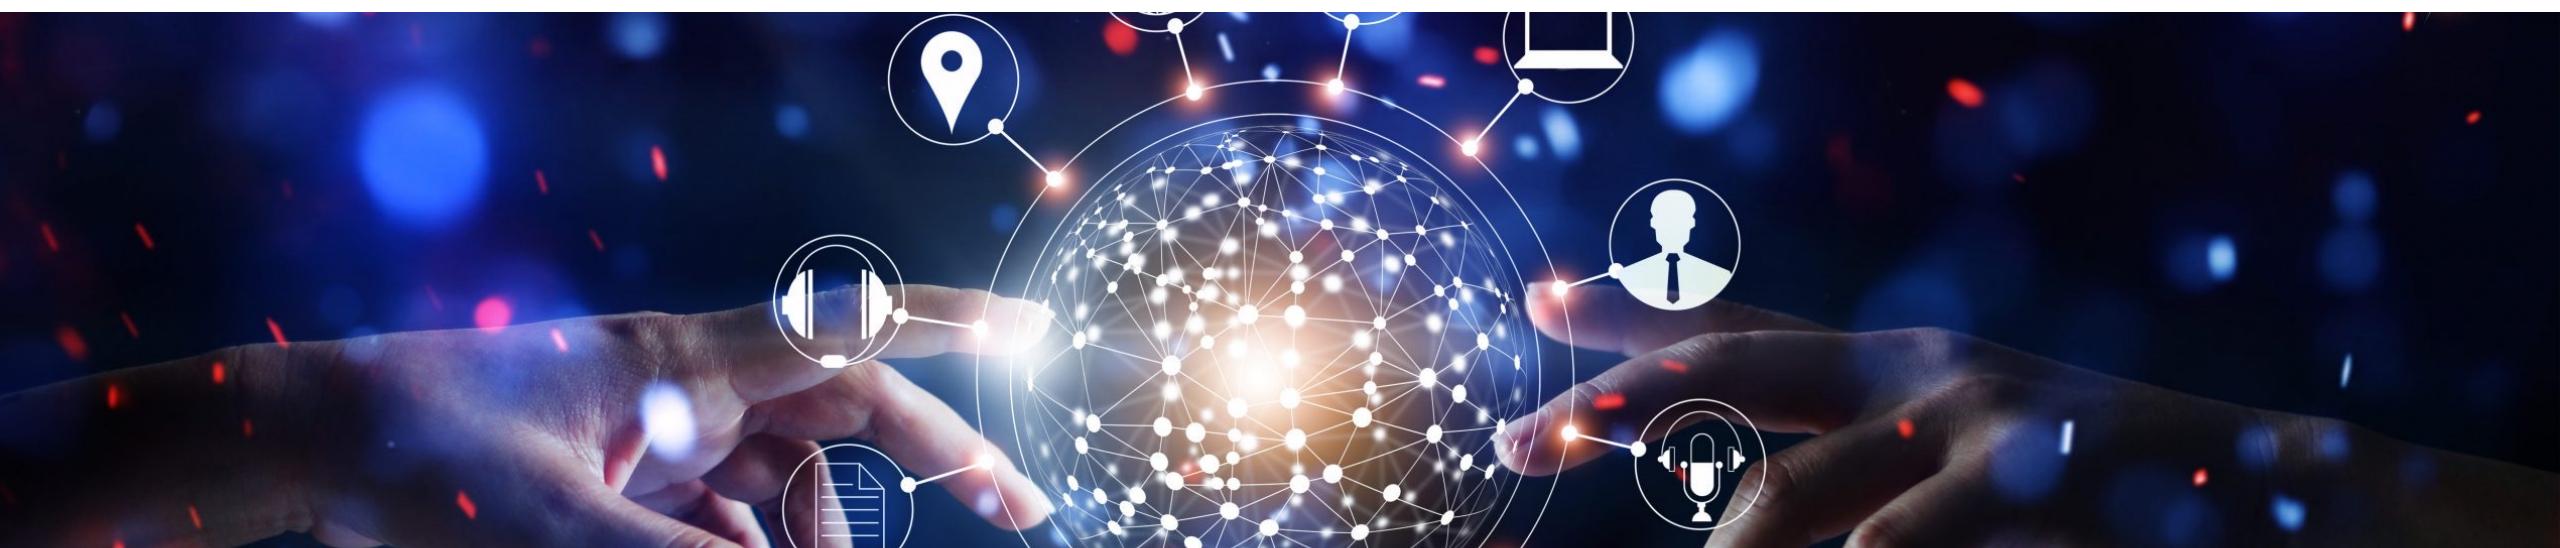


```
~$ sudo crictl inspect ef947d7856c12 | grep -iA 20 namespaces
"namespaces": [ {
    "type": "ipc",
    "path": "/proc/9665/ns/ipc"
},
{
    "type": "uts",
    "path": "/proc/9665/ns/uts"
},
{
    "type": "network",
    "path": "/proc/9665/ns/net"
},
]
user@worker-node:~$ ps -ef | grep 9665
root 9665 1 0 09:25 ? 00:00:00 /usr/bin/containerd-shim-kata-v2 -namespace
k8s.io -address /run/containerd/containerd.sock -publish-binary
/usr/bin/containerd -id a808b7b498fc088e228215fe4c4e6f74d40
```





Kubernetes pod security

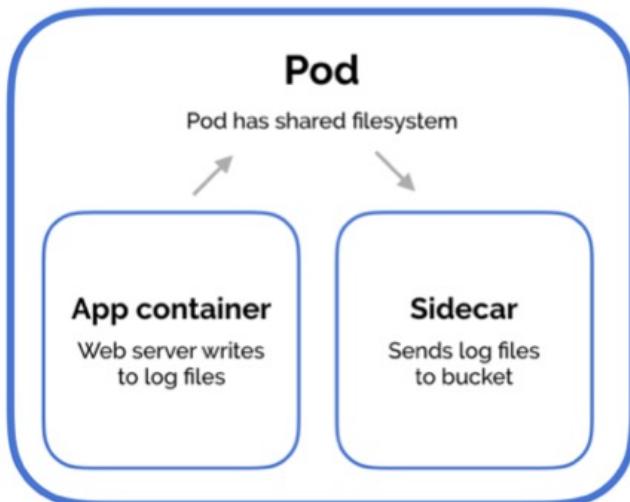


Cloud security principles

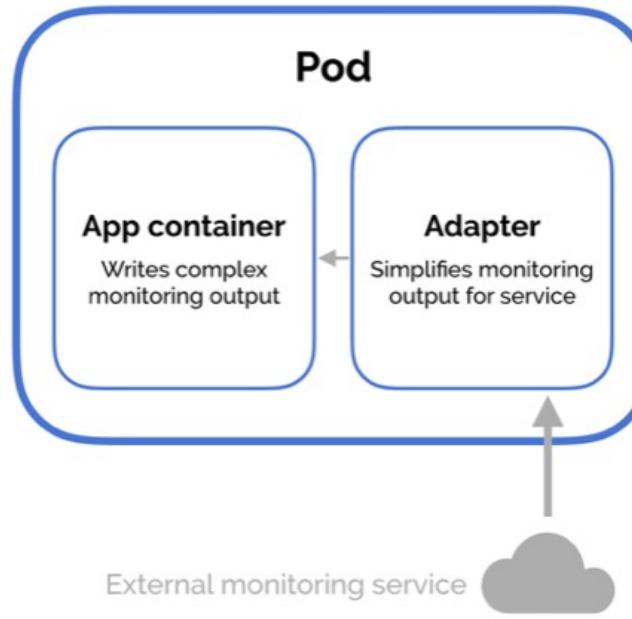
- Zero trust - AutheNtication
 - Who are you? (I'm Billy)
 - Everything must be authN
- Least privileges - AuthorizAtion
 - Hey Billy let me check what you can do
 - AuthZ for the fewest privileges
 - Use init containers running for short amount of time
- Anti patterns:
 - I don't need verification if it comes from inside

Multi container pod

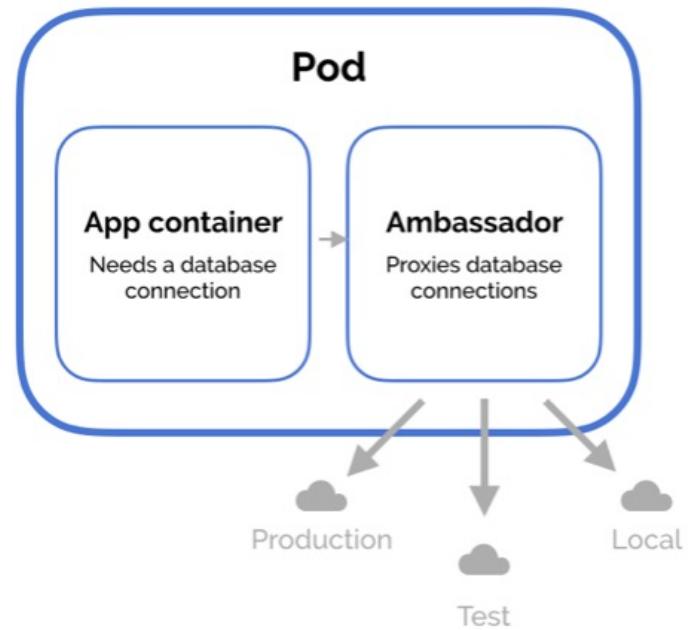
Sidecar



Adapter



Ambassador



Capabilities

```
root# grep Cap /proc/self/status
CapInh: 0000000000000000
CapPrm: 0000001fffffffff
CapEff: 0000001fffffffff
CapBnd: 0000001fffffffff
cncf$ grep Cap /proc/self/status
CapInh: 0000000000000000
CapPrm: 0000000000000000
CapEff: 0000000000000000
CapBnd: 0000001fffffffff
```

Capabilities

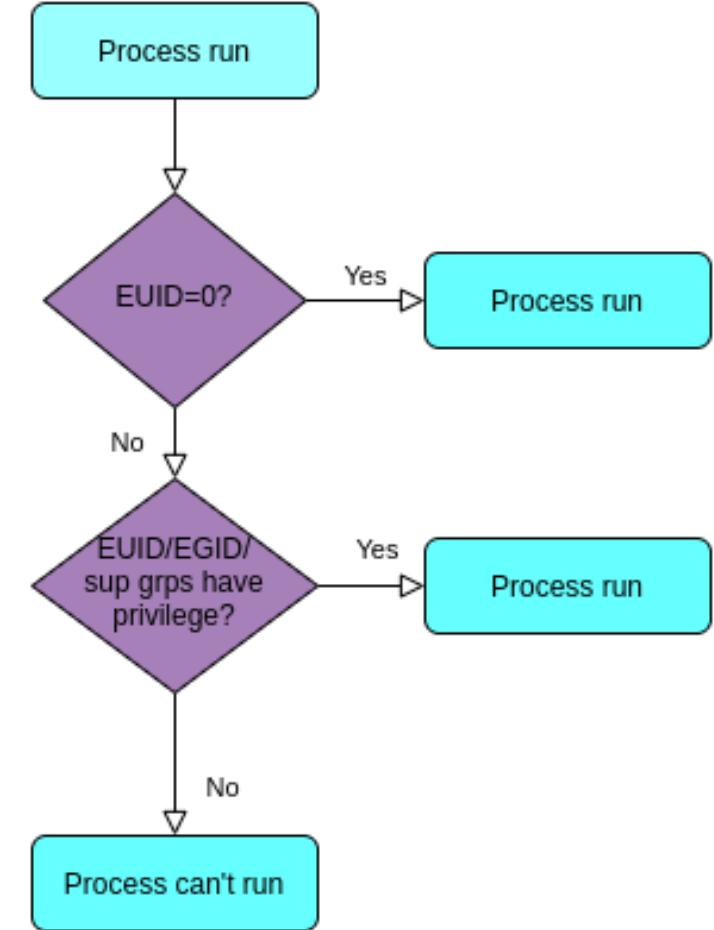


```
root# man capabilities
```

- privileged processes - user ID is 0
- unprivileged processes are subject to full permission checking based on the process's credentials

Example:

CAP_CHOWN
CAP_NET_ADMIN
CAP_SYS_TIME



More security methods

- SELinux - Linux kernel security module
 - Control access to the application
- AppArmor - Linux kernel security module
 - Applied to a Pod by specifying an AppArmor profile
- SecComp - security facility
 - Restrict System calls like a firewall



Kubernetes network security



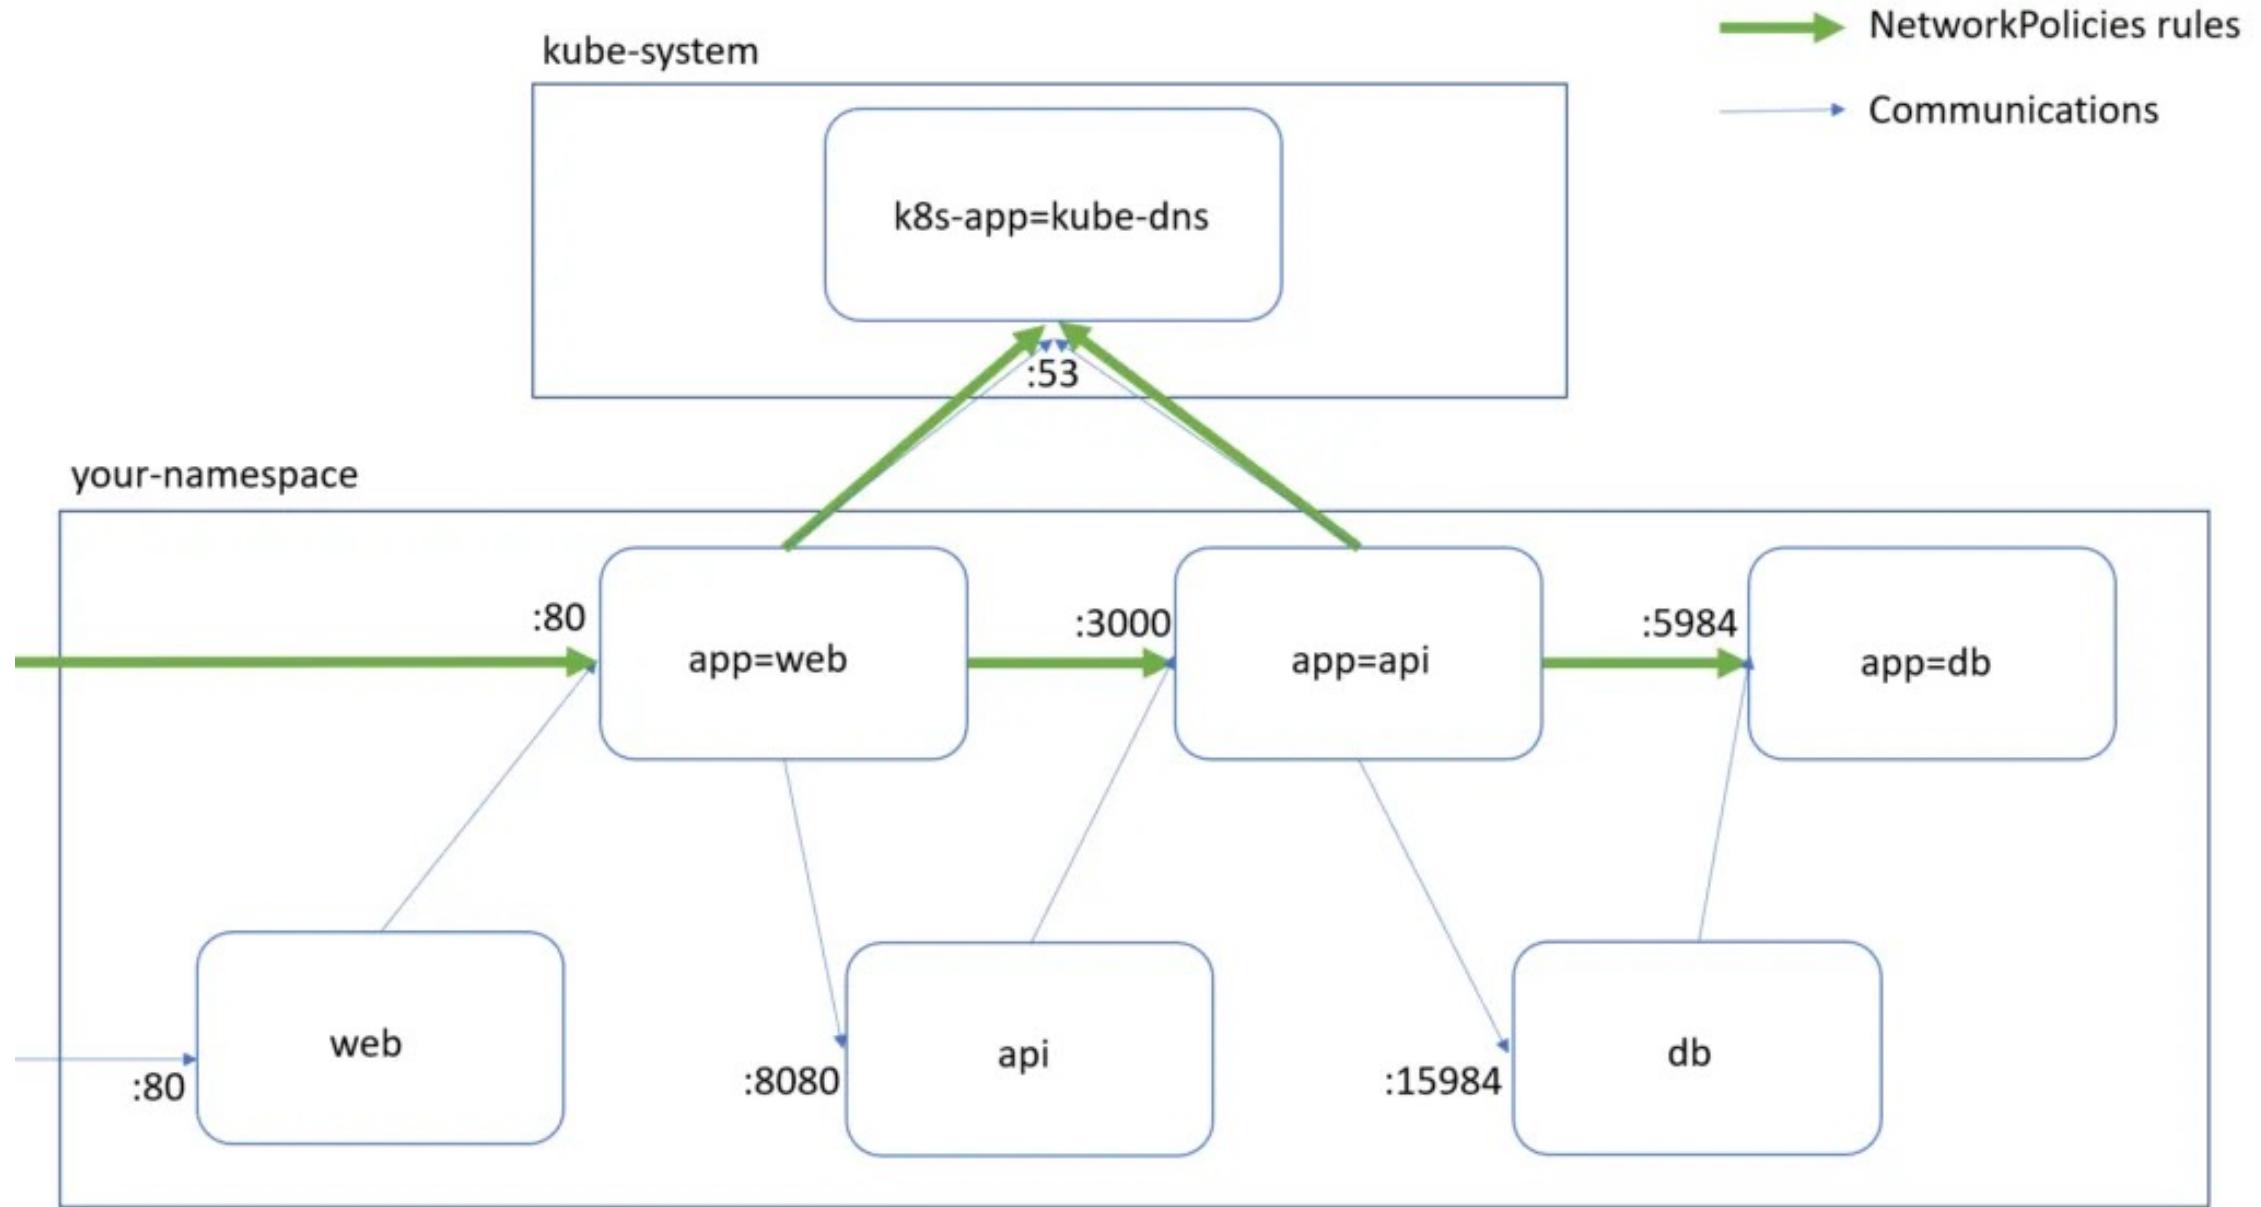
Network policies

Kubernetes Network Model

- Pods can communicate with all other pods on any other node *without NAT*
- Agents on a node can communicate with all pods on that node
- Control traffic flow at the IP address or port level (OSI layer 3 or 4)

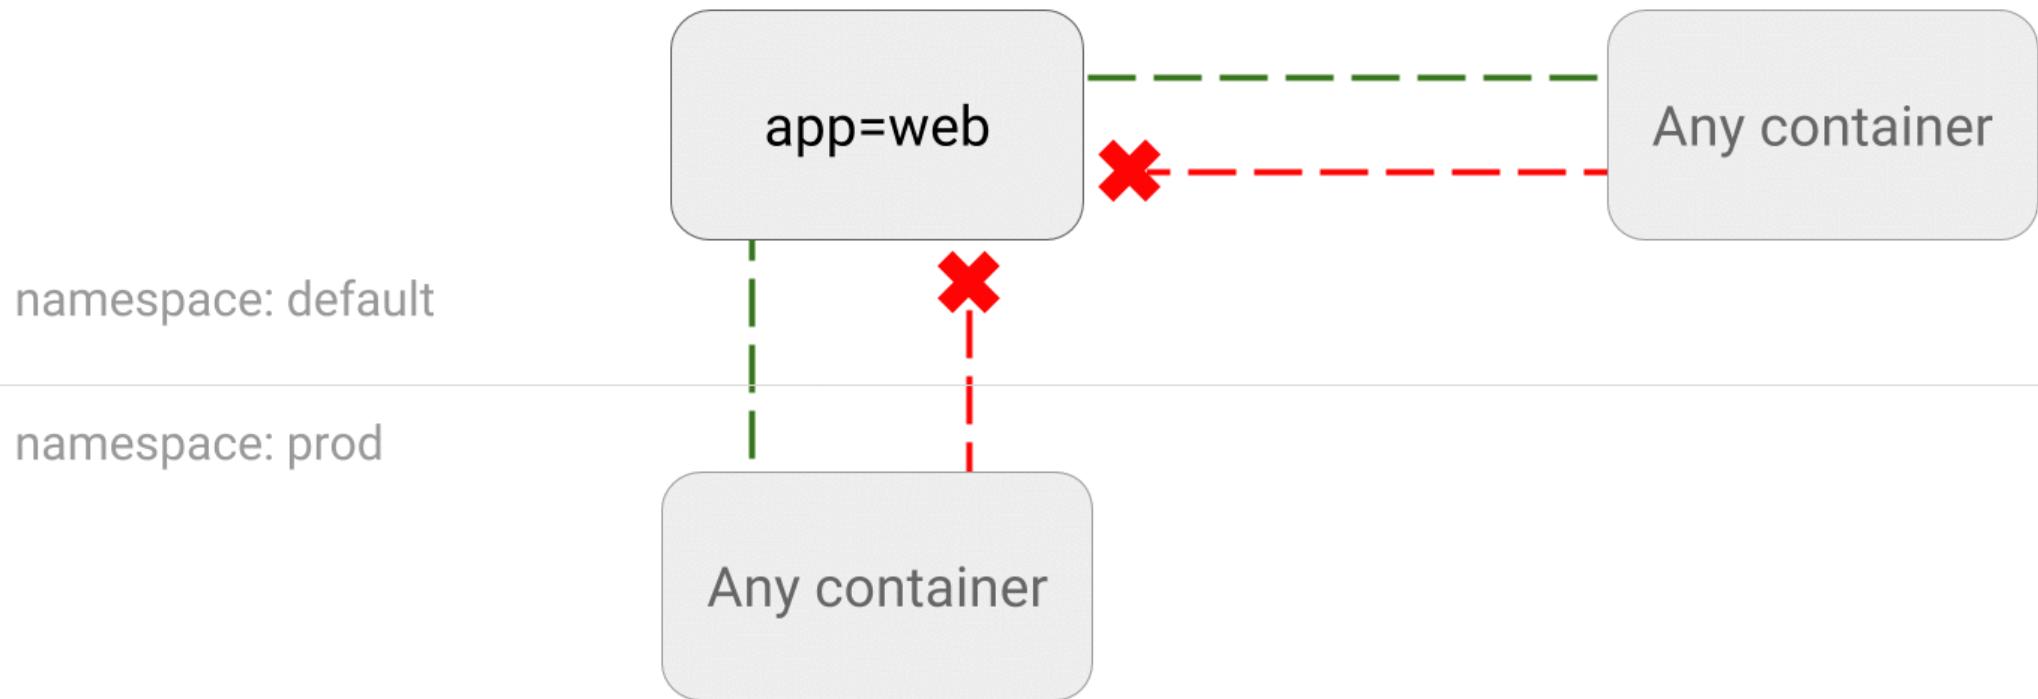
Network policies

- Container Network Interface Plugins:
 - Canal
 - Calico – popular
 - Weave Net – easy to install
 - Cilium – New generation CNI leverage eBPF
 - VMware NSX
 - Flannel – Do not support Network Policies



Network policies restrictions

Control Ingress and Egress traffic





Behavioral Analysis



Static Analysis

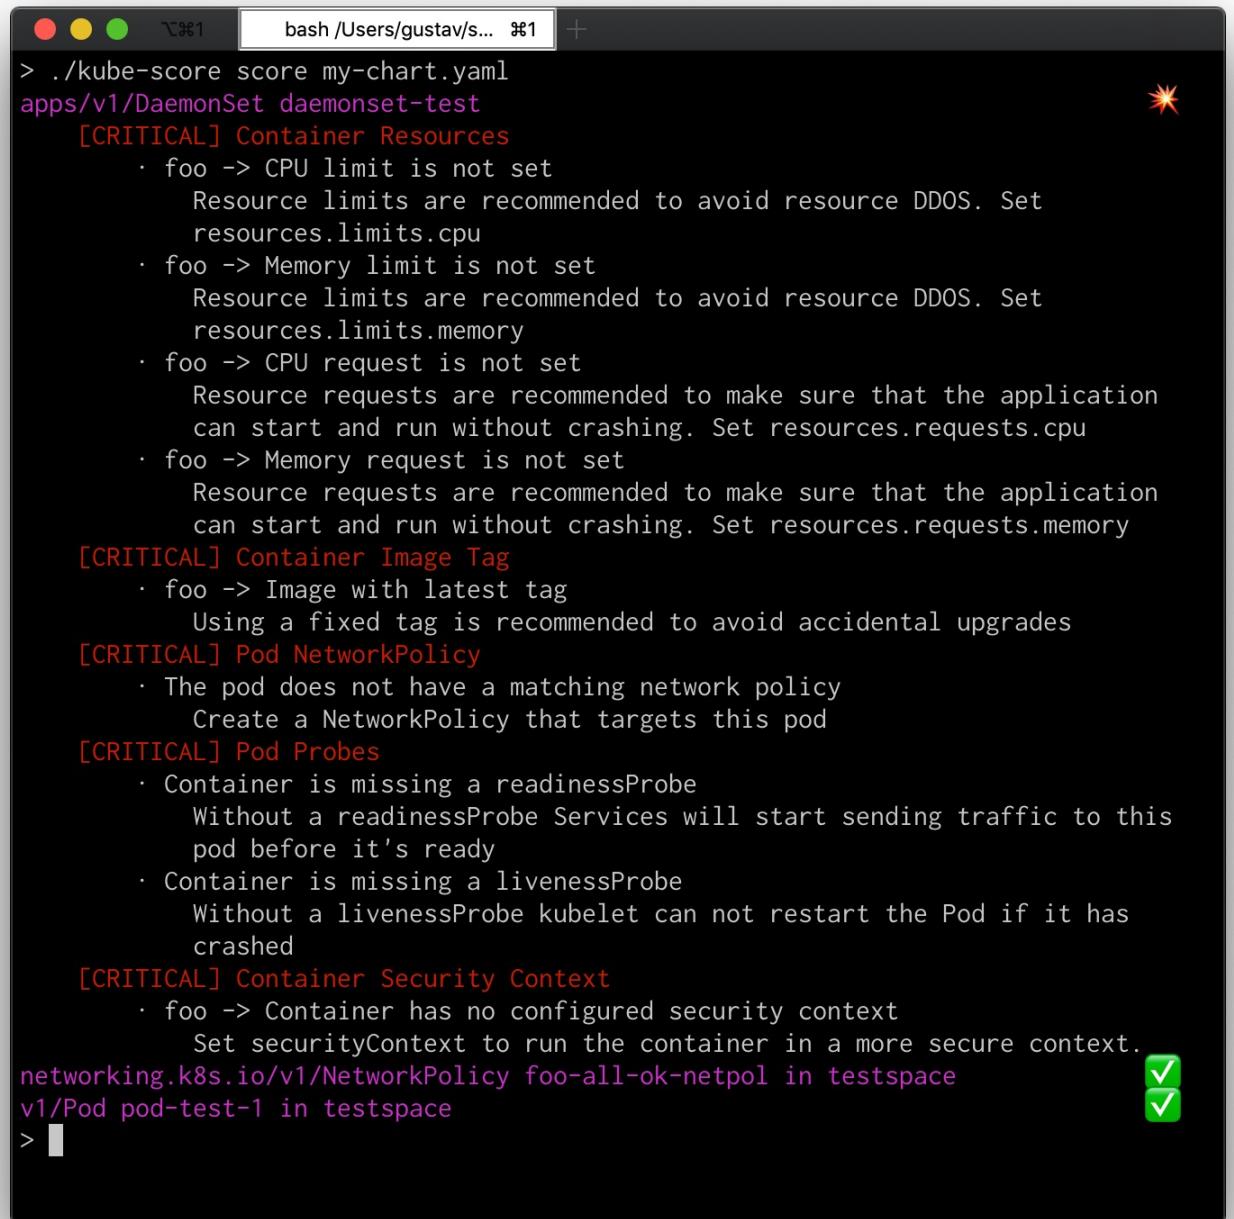
- Analyze your Infra as code
- Scan manifests for bad practices like:
 - Miss configuration
 - Escalation privileges
 - Hardcoded secrets (tokens, keys, creds...)
- Your computer editor > pre-commit > CICD
 - Linters - PyLint
 - Formatters - Black
 - Tools - pre-commit

- Kube-Score
- Implement it in your CICD
- Analyze manifests
 - Static YAML
 - Helm charts
 - Existing k8s resources

```

bash /Users/gustav/s... #1
> ./kube-score score my-chart.yaml
apps/v1/DaemonSet daemonset-test
[CRITICAL] Container Resources
  · foo -> CPU limit is not set
    Resource limits are recommended to avoid resource DDOS. Set
    resources.limits.cpu
  · foo -> Memory limit is not set
    Resource limits are recommended to avoid resource DDOS. Set
    resources.limits.memory
  · foo -> CPU request is not set
    Resource requests are recommended to make sure that the application
    can start and run without crashing. Set resources.requests.cpu
  · foo -> Memory request is not set
    Resource requests are recommended to make sure that the application
    can start and run without crashing. Set resources.requests.memory
[CRITICAL] Container Image Tag
  · foo -> Image with latest tag
    Using a fixed tag is recommended to avoid accidental upgrades
[CRITICAL] Pod NetworkPolicy
  · The pod does not have a matching network policy
    Create a NetworkPolicy that targets this pod
[CRITICAL] Pod Probes
  · Container is missing a readinessProbe
    Without a readinessProbe Services will start sending traffic to this
    pod before it's ready
  · Container is missing a livenessProbe
    Without a livenessProbe kubelet can not restart the Pod if it has
    crashed
[CRITICAL] Container Security Context
  · foo -> Container has no configured security context
    Set securityContext to run the container in a more secure context.
networking.k8s.io/v1/NetworkPolicy foo-all-ok-netpol in testspace
v1/Pod pod-test-1 in testspace
> █

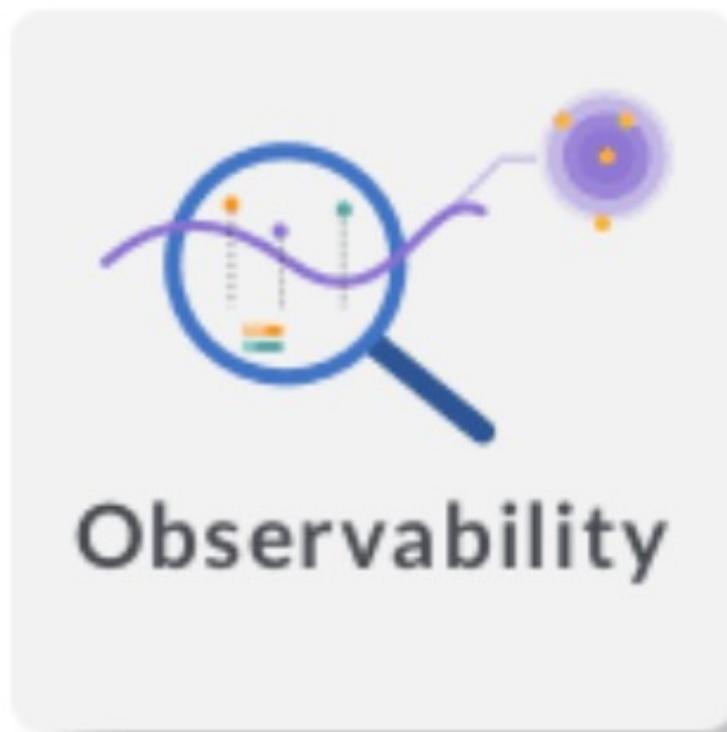
```



Running inside the cluster

- Do not touch it if it works
- I checked the image is **NOT** vulnerable
 - When ?
 - [1.021462] Kernel **panic()**
- Did you check your currently running images in your cluster ?
 - Remember security is **tight with time**





Observability usage

- Potential threat detection
 - Honeypots
- Behavior analysis
 - What an attacker is doing
- MITRE ATT & CK Matrix
 - Table of known attack techniques and behaviors documented by the enterprise security community

Observability usage

- STRIDE
 - Spoofing
 - Tampering
 - Repudiation
 - Information disclosure (privacy breach or data leak)
 - Denial of service
 - Elevation of privilege

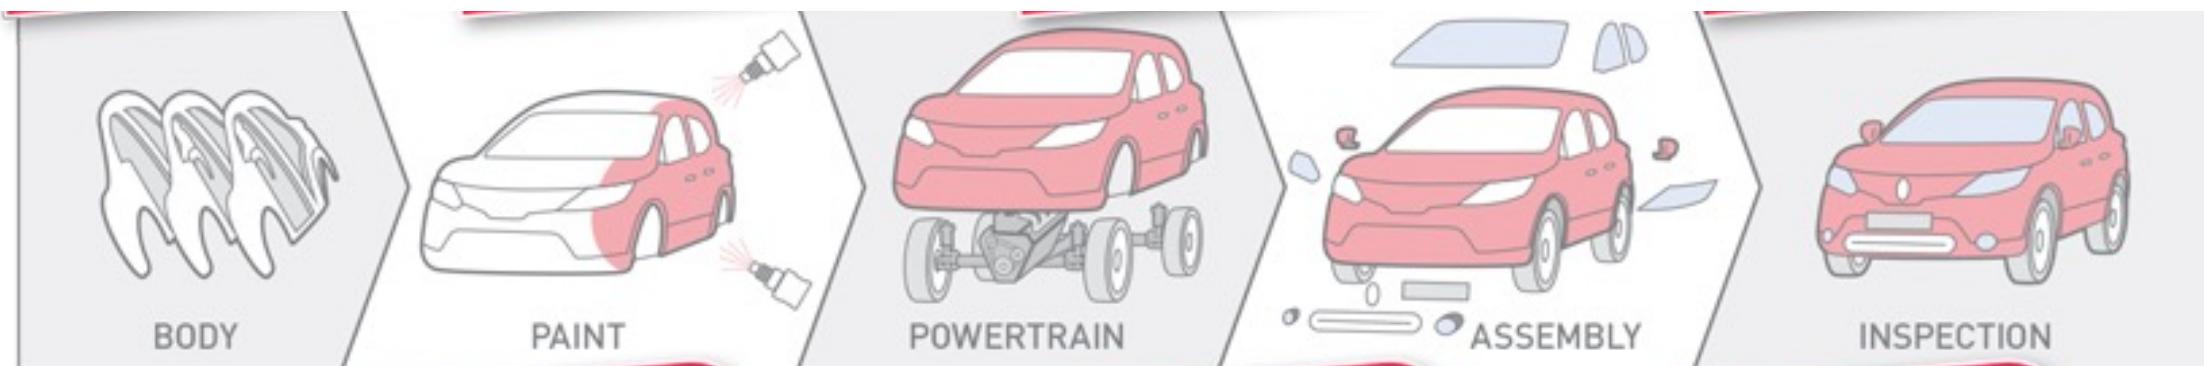
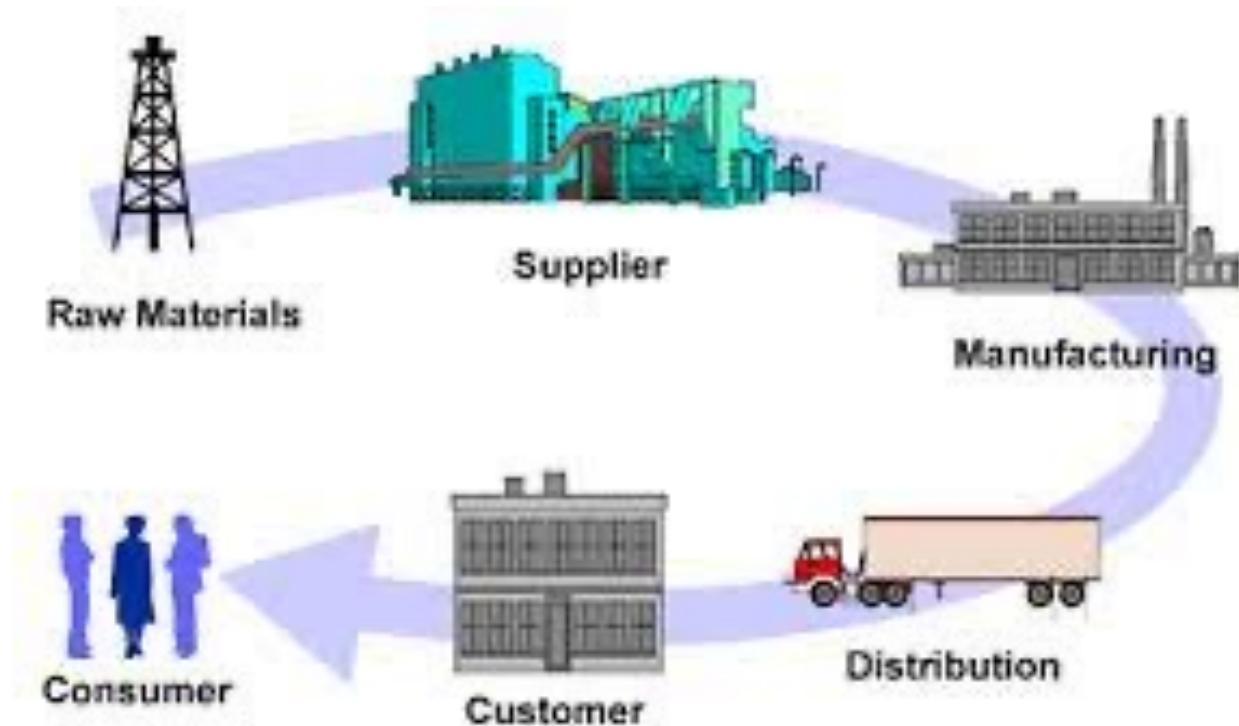


Supply Chain Security



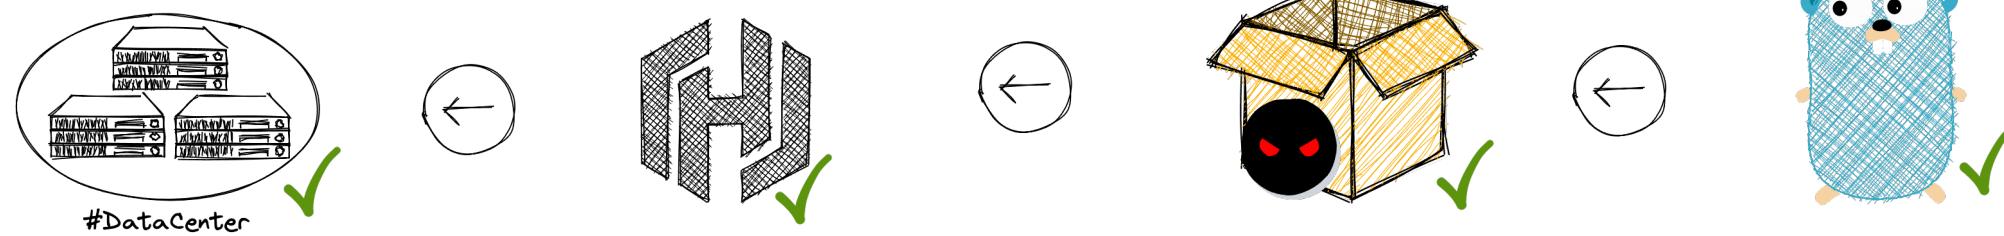
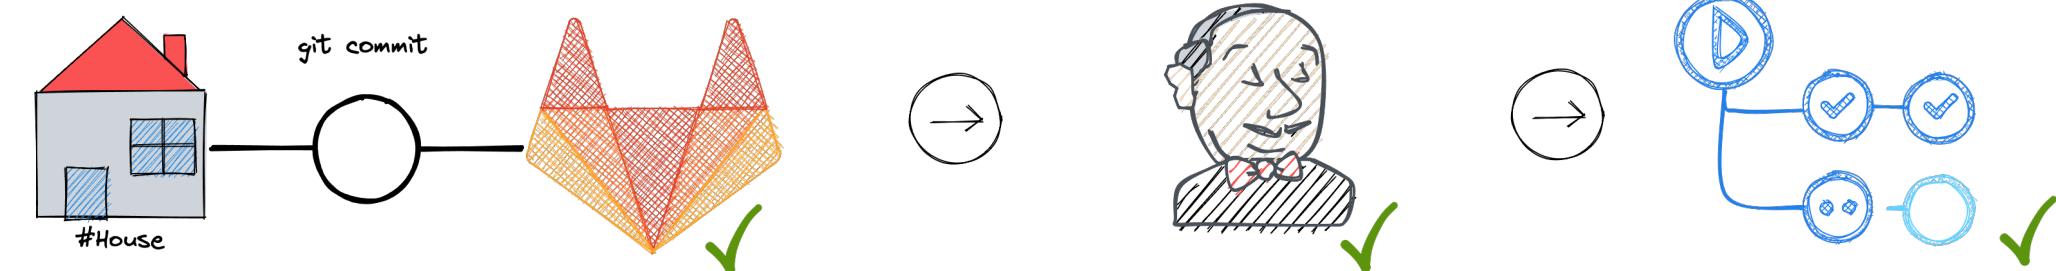
Supply chain

What is actually a supply chain ?

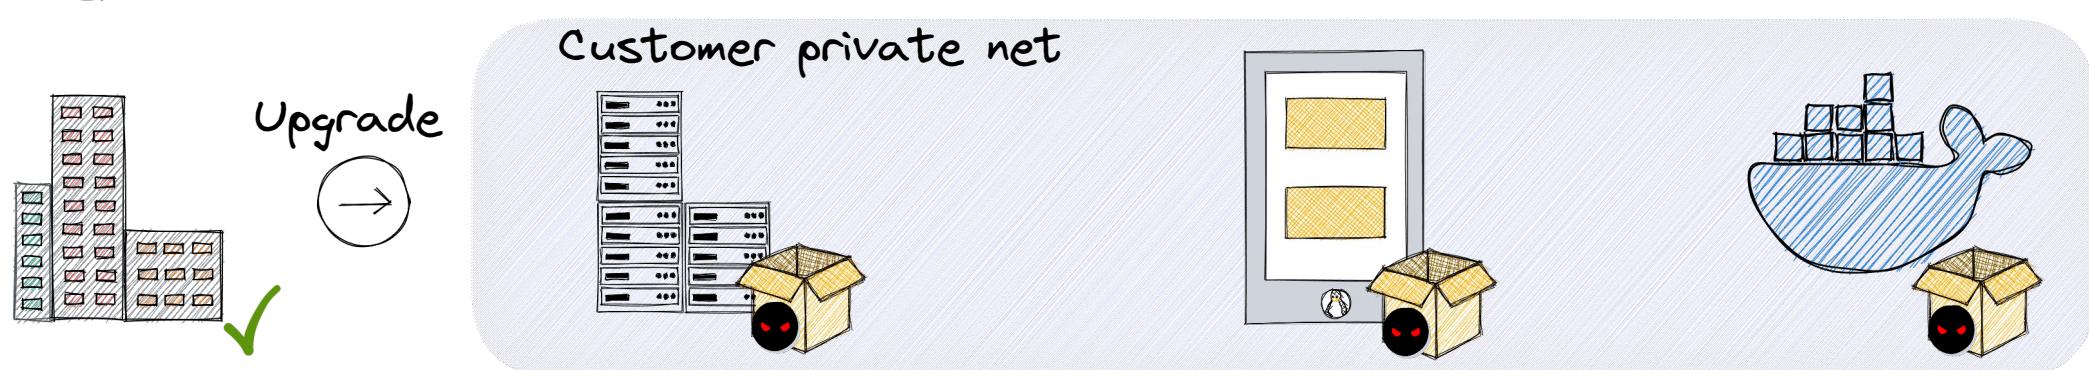


Supply chain tools

- You use 'Trivy' to scan your images
 - Guest what we hacked 'Tviry'
 - Guest what now our malware is inside the image and you don't know because 'Tviry' can't catch it
 - Use multiple scanners bypass if 3 of 5 are meeting your policies



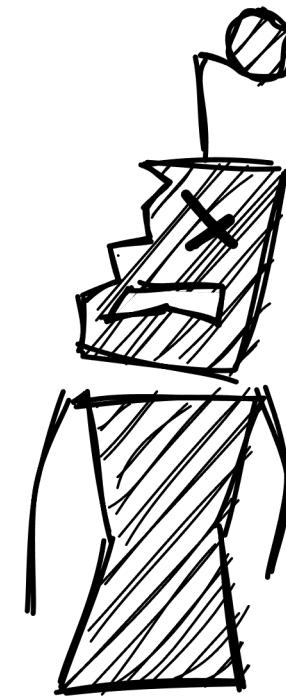
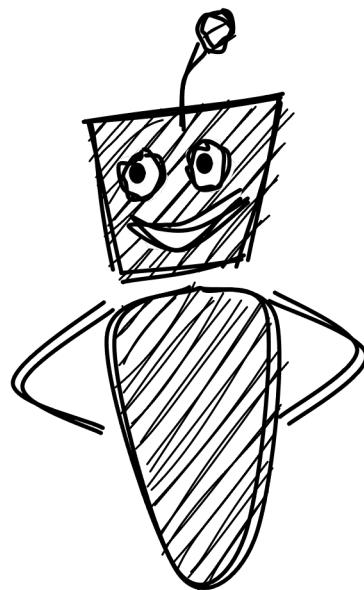
Download update



It's not a job it's a
passion

Q&A

01000001 01101110 01111001 00100000
01110001 01110101 01100101 01110011
01110100 01101001 01101111 01101110 01110011
00111111



Sources

- Kata containers: <https://thenewstack.io/the-road-to-kata-containers-2-0/>
- Alpine malware miner - <https://www.trendmicro.com/vinfo/fr/security/news/virtualization-and-cloud/malicious-docker-hub-container-images-cryptocurrency-mining>
- Monero mining: <https://unit42.paloaltonetworks.com/cryptojacking-docker-images-for-mining-monero/>
- Google 6 layers of security
- <https://www.youtube.com/watch?v=kd33UVZhnAA>
- sidecar: <https://matthewpalmer.net/kubernetes-app-developer/articles/multi-container-pod-design-patterns.html>
- network policies: <https://cloudogu.com/en/blog/k8s-app-ops-part-1>, <https://github.com/ahmetb/kubernetes-network-policy-recipes> ,
<https://cloudblogs.microsoft.com/opensource/2019/10/17/tutorial-calico-network-policies-with-azure-kubernetes-service/>
- CNI plugins: <https://github.com/containernetworking/cni#3rd-party-plugins>
- Observability: https://linkedin.github.io/school-of-sre/level101/metrics_and_monitoring/observability/